

Parallel Programming Exercises

Andreas Hollmann

Technische Universität München

April 23, 2013

Organization

- ▶ Lecture could start at:
 - ▶ 12:00
 - ▶ 12:15
 - ▶ 12:30
- ▶ Poll is available at <http://doodle.com/uzy7zfg6gaatwgkg>
- ▶ Duration: as long as we need, up to 90 min.
- ▶ Additional assistance on assignments:
 - ▶ Tuesday 10:00 - 12:00 room 01.04.011
 - ▶ Wednesday 14:00 - 16:00 room 01.04.011
- ▶ Detailed discussion of programming techniques and assignments in this lecture
- ▶ My email address is: hollmann@in.tum.de
- ▶ Website will be updated and extended in the next days

Assignments

- ▶ Neither graded nor corrected
- ▶ Still very important
- ▶ Small programming tasks in final exam
 - ▶ without practice you will run out of time or do it wrong
- ▶ Solutions will be made public
- ▶ Topics
 - ▶ Pthreads (Posix Threads)
 - ▶ OpenMP (Open Multi-Processing)
 - ▶ MPI (Message Passing Interface)
- ▶ Code examples are in C99
- ▶ C++ won't be covered

Additional Book Resources

- ▶ Students at TUM have access to Safari Books Online
 - ▶ <http://proquest.tech.safaribooksonline.de.eaccess.ub.tum.de>
 - ▶ Login: user@mytum.de or user@tum.de and your password
 - ▶ Online access to really good books
 - ▶ Printing single chapters is possible
 - ▶ Searching for keywords across all books
- ▶ Recommended Books
 - ▶ An Introduction to Parallel Programming, by Peter Pacheco
 - ▶ Programming with Posix Threads, by David Butenhof
 - ▶ The Linux Programming Interface, by Michael Kerrisk
 - ▶ Patterns for Parallel Programming, by Timothy G. Mattson; Beverly A. Sanders; Berna L. Massingill


Additional Video Resources

- ▶ ParLAB
 - ▶ Very good presentations on different parallel programming topics at University of California, Berkeley
- ▶ Ulrich Drepper
 - ▶ Scalable Parallel Programming Techniques
 - ▶ Why knowing your hardware is important

Course Prerequisites

- ▶ Knowledge of C
 - ▶ What is a global or static variable?
 - ▶ What does const mean in C? Is `'const int *c = &b;` and `int const *c = &b;` the same?
- ▶ C books
 - ▶ (C89) The C Programming Language, Second Edition, by Brian W. Kernighan; Dennis M. Ritchie
 - ▶ (C99) C Primer Plus, Fifth Edition, by Stephen Prata
- ▶ Experience with Linux Command Line
- ▶ Resources
 - ▶ Book: The Linux Command Line
 - ▶ Basic video introduction: The Shell
- ▶ Knowing GCC
 - ▶ An Introduction to GCC, by Brian Gough

35 Years of Microprocessor Trend Data



Original data collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond and C. Batten
Dotted line extrapolations by C. Moore

Year 2005: The Free Lunch Is Over

- ▶ A Fundamental Turn Toward Concurrency in Software
- ▶ Software doesn't get (much) faster with the next microprocessor generation
- ▶ Prior to 2004, single-threaded floating-point performance climbed, at 64% per year: a doubling period of 73 weeks. After that, it leveled off at 21% per year.
- ▶ Software developers have to rewrite their applications to use multiple processors in order to speed them up
- ▶ Parallel Programming is hard
 - ▶ to write - complex APIs and needs more code than serial version
 - ▶ to do it correctly - it's easy to introduce new bugs
 - ▶ to debug, order of thread execution is undefined
 - ▶ to make it scalable - will your applications scale with more cores?
 - ▶ better qualified developers are necessary

Is parallel programming that complex?

- ▶ It depends on the problem
- ▶ For completely independent tasks it's rather easy
 - ▶ Embarrassingly Parallel - parameter studies or serving static files on a webserver to multiple users
- ▶ It gets more and really complex with dependencies between tasks and shared data
- ▶ Some problems are simple to understand, but really hard to implement efficiently, e. g. Dijkstra's shortest-path algorithm.
- ▶ Scientific Algorithms are ok
 - ▶ Often ends up in some stencil or matrix computation

Posix Thread Programming

- ▶ Single process with multiple paths of execution running in parallel
 - ▶ Functions are started in new threads and will run concurrently
- ▶ Usage:
 - ▶ In GUIs: one thread updates the GUI, second thread does actual work
 - ▶ To speed up IO intensive workloads, one thread does IO, other threads do computation
 - ▶ To speed up computations, by using all CPUs
- ▶ Posix Threads were defined in 1995 (IEEE Std 1003.1c-1995)
- ▶ Threads are light-weight compared to processes
- ▶ All threads within one process share
 - ▶ Memory resources (data, heap, code area)
 - ▶ File descriptors
 - ▶ Environment (privileges, working directory)
- ▶ Each thread has a private stack

Why are threads faster than processes?

- ▶ Creating a new process with `fork()` has a big overhead: whole memory must be copied
 - ▶ Waste of memory resources!
- ▶ Synchronization with processes usually involves system calls.

Output ps tool

```
hollmann@lxhalle:~$ ps -fLmU hollmann
```

UID	PID	LWP	C	NLWP	CMD
hollmann	16233	—	0	1	sshd: hollmann@pts/29
hollmann	—	16233	0	—	—
hollmann	16234	—	0	1	—bash
hollmann	—	16234	0	—	—
hollmann	16279	—	0	1	sshd: hollmann@pts/31
hollmann	—	16279	0	—	—
hollmann	16280	—	0	1	—bash
hollmann	—	16280	0	—	—
hollmann	28347	—	99	2	./parallel_app
hollmann	—	28347	81	—	—
hollmann	—	28348	74	—	—
hollmann	28349	—	0	1	ps -fLmU hollmann
hollmann	—	28349	0	—	—

Posix Threads on Linux

- ▶ Pthreads are mapped 1:1 to Linux threads
- ▶ A process can have many threads (at least one) 1:N and is identified by its process id (PID)
- ▶ Each linux thread has an unique thread id (TID) and belongs to exactly one process
- ▶ There is no defined printable id of a pthread, on Linux it is possible to use TID

Hello World Example

```
1  int main()  
2  {  
3      printf("Hello World!\n");  
4  }
```

Hello World with Pthreads Ver. 0

```
1 void * hello()  
2 {  
3     printf("Hello World from pthread!\n");  
4     return NULL;  
5 }  
6  
7 void main()  
8 {  
9     pthread_t thread;  
10  
11     pthread_create(&thread, NULL, &hello, NULL);  
12  
13     printf("Hello World from main!\n");  
14  
15     pthread_join(thread, NULL);  
16 }
```

Create Pthreads

```
1  int pthread_create(pthread_t *thread,  
2                          const pthread_attr_t *attr,  
3                          void *(*start_routine) (void *),  
4                          void *arg);
```

- ▶ pthread_t *thread,
 - ▶ Pointer to thread identifier.
- ▶ const pthread_attr_t *attr
 - ▶ Optional pointer to pthread_attr_t to define behavior, if NULL defaults are used.
- ▶ void *(*start_routine) (void *),
 - ▶ Pointer to function prototype that is started. Function takes void pointer as argument and returns a void pointer.
- ▶ void *arg
 - ▶ Pointer to the argument that is used for the executed function.

Waiting for Pthread to finish

```
1  int pthread_join(pthread_t thread,  
2                      void **retval);
```

- ▶ pthread_t thread,
 - ▶ Pointer to thread identifier, for which this function is waiting.
- ▶ void **retval
 - ▶ Optional pointer pointing to a void pointer. This can be used to return data of undefined size.

Compile & Output

```
gcc --std=gnu99 -pthread -Wall  
    -o hello_world hello_world.c
```

```
Hello World from main!
```

```
Hello World from pthread!
```

Hello World with Pthreads Ver. 1

```
1 void main()  
2 {  
3     long num_threads = 3; pthread_t *thread;  
4  
5     thread = malloc((num_threads * sizeof(*thread)));  
6  
7     for (int i = 0; i < num_threads; i++)  
8         pthread_create(thread + i, NULL, &hello, NULL );  
9  
10    for (int i = 0; i < num_threads; i++)  
11        pthread_join(thread[i], NULL );  
12 }
```

Output

```
[user]$ ./hello_world_thread_1  
Hello World from pthread!  
Hello World from pthread!  
Hello World from pthread!
```

Hello World with Pthreads Ver. 2

```
1 void * hello(void *ptr)
2 {
3     printf("Hello World from pthread %d!\n", *((int*)ptr));
4     return NULL ;
5 }
```

Hello World with Pthreads Ver. 2

```
1  void main()
2  {
3      int num_threads = 3;  pthread_t *thread;  int *thread_arg;
4
5      thread = malloc(num_threads * sizeof(*thread));
6      thread_arg = malloc(num_threads * sizeof(*thread_arg));
7
8      for (int i = 0; i < num_threads; i++)
9      {
10         thread_arg[i] = i;
11         pthread_create(thread + i, NULL, &hello, thread_arg + i);
12     }
13
14     for (int i = 0; i < num_threads; i++)
15         pthread_join(thread[i], NULL );
16 }
```

Output

```
[user]$ ./hello_world_thread_2  
Hello World from thread 0!  
Hello World from thread 1!  
Hello World from thread 2!
```

Hello World with Pthreads Ver. 3

```
1
2 pid_t gettid() { return (pid_t) syscall( __NR_gettid ); }
3
4 struct pthread_args
5 {
6     long thread_id;
7     long num_threads;
8 };
9
10 void * hello(void *ptr) {
11     struct pthread_args *arg = ptr;
12     printf("Hello World from pthread %ld of %ld PID = %d TID = %d\n",
13           arg->thread_id,
14           arg->num_threads,
15           getpid(),
16           gettid());
17
18     return NULL ;
19 }
```


Hello World with Pthreads Ver. 3

```
1  void main()
2  {
3      long num_threads = 3;
4      pthread_t *thread;
5      struct pthread_args *thread_arg;
6
7      thread = malloc(num_threads * sizeof(*thread));
8      thread_arg = malloc(num_threads * sizeof(*thread_arg));
9
10     for (int i = 0; i < num_threads; i++)
11     {
12         thread_arg[i].thread_id = i;
13         thread_arg[i].num_threads = num_threads;
14         pthread_create(thread + i, NULL, &hello_pthread, thread_arg + i);
15     }
16
17     for (int i = 0; i < num_threads; i++)
18         pthread_join(thread[i], NULL );
19 }
```

Output

```
[user]$ ./hello_world_thread_3  
Hello World from pthread 1 of 3 PID = 23750 TID = 23752!  
Hello World from pthread 0 of 3 PID = 23750 TID = 23751!  
Hello World from pthread 2 of 3 PID = 23750 TID = 23753!
```

Assignment: Numerical integration of PI

$$\pi = 4 \cdot \int_0^1 \sqrt{1-x^2} \, dx$$

Assignment: Parallelize this Code

```
1  #include <stdio.h>
2  #include <math.h>
3  #define STEPS 1000000
4
5  void main()
6  {
7      double step_size = 1.0/STEPS, t = 0.5 * step_size, sum = 0;
8
9      while(t < 1.0)
10     {
11         sum += sqrt(1-t*t) * step_size;
12         t += step_size;
13     }
14     sum *= 4;
15
16     printf("Computed PI = %.10lf\n", sum);
17     printf("Difference to Reference is %.10lf\n", M_PI - sum);
18 }
```