

Parallel Programming Exercises

Andreas Hollmann

Technische Universität München

April 23, 2013

Organization

- ▶ This course will start at 12:00
- ▶ No groups, no lab courses
- ▶ Assistance sessions as announced last time: on Tuesdays between 10:00 - 12:00 and on Wednesdays from 14:00 - 16:00, starting next week

Assignment: Parallelize this Code

```
1  #include <stdio.h>
2  #include <math.h>
3  #define STEPS 1000000
4
5  void main()
6  {
7      double step_size = 1.0/STEPS, t = 0.5 * step_size, sum = 0;
8
9      while(t < 1.0)
10     {
11         sum += sqrt(1-t*t) * step_size;
12         t += step_size;
13     }
14     sum *= 4;
15
16     printf("Computed PI = %.10lf\n", sum);
17     printf("Difference to Reference is %.10lf\n", M_PI - sum);
18 }
```

One PI Solution 1/4

```
1  #define STEPS 1000000
2  #define STEP_SIZE 1.0/STEPS
3  #define THREADS 3
4
5  struct pthread_args
6  {
7      double lower;
8      double upper;
9      double local_sum;
10 };
```

One PI Solution 2/4

```
1 void * pi_thread(void *ptr)
2 {
3     double low = 0.5 * STEP_SIZE +
4             ((struct pthread_args*)ptr)->lower;
5     double upp = ((struct pthread_args*)ptr)->upper;
6     double tsum = 0;
7
8     while(low < upp)
9     {
10         tsum += sqrt(1-low*low) * STEP_SIZE;
11         low += STEP_SIZE;
12     }
13     ((struct pthread_args*)ptr)->local_sum = tsum;
14
15     return NULL;
16 }
```

One PI Solution 3/4

```
1  void main()
2  {
3      long num_threads = 10000; double sum = 0;
4      pthread_t *thread; struct pthread_args *thread_arg;
5      thread = malloc(num_threads * sizeof(*thread));
6      thread_arg = malloc(num_threads * sizeof(*thread_arg));
7
8      for (int i = 0; i < num_threads; i++)
9      {
10         thread_arg[i].lower = (i+0) * (1.0/((double)num_threads));
11         thread_arg[i].upper = (i+1) * (1.0/((double)num_threads));
12         pthread_create(thread+i, NULL, &pi_thread, thread_arg+i);
13     }
14     for (int i = 0; i < num_threads; i++)
15     {
16         pthread_join(thread[i], NULL );
17         sum += 4 * thread_arg[i].local_sum;
18     }
19 }
```

One PI Solution 3/4

```
1 void main()  
2 {  
3     ...  
4     printf("Reference PI = %.10lf Computed PI = %.10lf\n", M_PI,  
5     printf("Difference to Reference is %.10lf\n", M_PI - sum);  
6 }
```

Output PI

```
$ time ./pi
```

```
Reference PI = 3.1415926536 Computed PI = 3.1415926539
```

```
Difference to Reference is -0.0000000003
```

```
real    0m0.021s
```

```
user    0m0.020s
```

```
sys     0m0.000s
```

```
$ time ./pi_thread_0
```

```
Reference PI = 3.1415926536 Computed PI = 3.1415926539
```

```
Difference to Reference is -0.0000000003
```

```
real    0m0.487s
```

```
user    0m0.127s
```

```
sys     0m0.660s
```


Why knowing your hardware is important?

- ▶ In order to set the right number of threads.
- ▶ To run the threads on the right CPUs

CPU Information 1/4

```
$ cat /proc/cpuinfo
```

```
processor          : 0
```

```
vendor_id         : GenuineIntel
```

```
cpu family        : 6
```

```
model             : 37
```

```
model name        : Intel(R) Core(TM) i7 CPU M 620  @ 2.670
```

```
stepping          : 2
```

```
microcode         : 0xd
```

```
cpu MHz           : 1199.000
```

```
cache size        : 4096 KB
```

```
physical id       : 0
```

```
siblings          : 4
```

```
core id           : 0
```

```
cpu cores         : 2
```

```
apicid            : 0
```

CPU Information 2/4

```
processor      : 1
vendor_id     : GenuineIntel
cpu family    : 6
model         : 37
model name    : Intel(R) Core(TM) i7 CPU M 620  @ 2.67GHz
stepping      : 2
microcode     : 0xd
cpu MHz       : 1199.000
cache size    : 4096 KB
physical id   : 0
siblings      : 4
core id       : 2
cpu cores     : 2
apicid        : 4
```

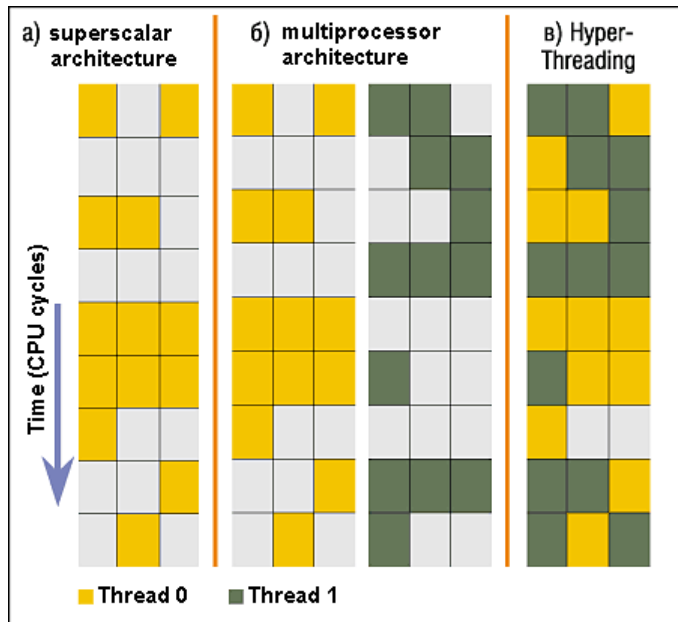
CPU Information 3/4

```
processor      : 2
vendor_id     : GenuineIntel
cpu family    : 6
model         : 37
model name    : Intel(R) Core(TM) i7 CPU M 620  @ 2.67GHz
stepping      : 2
microcode     : 0xd
cpu MHz       : 1199.000
cache size    : 4096 KB
physical id   : 0
siblings      : 4
core id       : 0
cpu cores     : 2
apicid        : 1
```

CPU Information 4/4

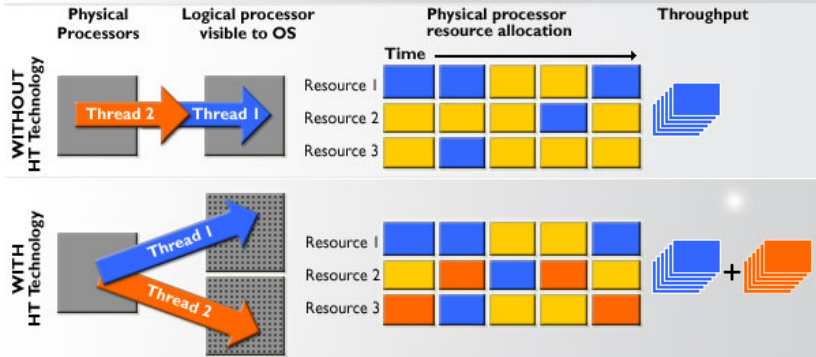
```
processor      : 3
vendor_id     : GenuineIntel
cpu family    : 6
model         : 37
model name    : Intel(R) Core(TM) i7 CPU M 620  @ 2.67GHz
stepping      : 2
microcode     : 0xd
cpu MHz       : 1199.000
cache size    : 4096 KB
physical id   : 0
siblings      : 4
core id       : 2
cpu cores     : 2
apicid        : 5
```

Hyperthreading (SMT)



Hyperthreading (SMT)

How Hyper-Threading Technology Works



Hyperthreading (SMT)

- ▶ Superscalar architectures have several functional units that can work in parallel
- ▶ Functional units are not well utilized with only one thread (stream of instruction)
- ▶ SMT duplicates part of a real CPU core, but shares functional units with other threads
- ▶ Improves the utilization of the functional units for certain applications or combinations of applications
 - ▶ One application does integer computation the other does floating point calculation

PI Measurements with 4 and 2 Threads

```
$ time ./pi_pthread_0 \begin
```

```
Reference PI = 3.1415926536 Computed PI = 3.1415926623
```

```
Difference to Reference is -0.0000000087 <-- compare
```

```
real    0m8.865s
```

```
user    0m34.323s <-- 4 threads
```

```
sys     0m0.027s
```

```
$ time ./pi_pthread_0
```

```
Reference PI = 3.1415926536 Computed PI = 3.1415926626
```

```
Difference to Reference is -0.0000000090 <-- differs
```

```
real    0m8.867s
```

```
user    0m17.673s <-- 2 threads
```

```
sys     0m0.000s
```

Limiting the CPU_SET with taskset command

```
$ time taskset -c 0 ./pi_pthread_0
```

```
Reference PI = 3.1415926536 Computed PI = 3.1415926626
```

```
Difference to Reference is -0.0000000090
```

```
real    0m16.568s
```

```
user    0m16.513s
```

```
sys     0m0.007s
```

```
$ time taskset -c 0,2 ./pi_pthread_0
```

```
Reference PI = 3.1415926536 Computed PI = 3.1415926626
```

```
Difference to Reference is -0.0000000090
```

```
real    0m16.577s
```

```
user    0m33.033s
```

```
sys     0m0.017s
```

Limiting the CPU_SET with taskset command

```
$ time taskset -c 0,2 ./pi_pthread_0
```

```
Reference PI = 3.1415926536 Computed PI = 3.1415926626
```

```
Difference to Reference is -0.0000000090
```

```
real    0m16.577s
```

```
user    0m33.033s
```

```
sys     0m0.017s
```

```
$ time taskset -c 0,1 ./pi_pthread_0
```

```
Reference PI = 3.1415926536 Computed PI = 3.1415926626
```

```
Difference to Reference is -0.0000000090
```

```
real    0m8.876s
```

```
user    0m17.687s
```

```
sys     0m0.000s
```

Linux 2.6.23 (CFS) Scheduling Overview

- ▶ One run queue per (logical) CPU
- ▶ Active threads are placed in one of these queues
- ▶ Thread runs until his time slice is over or it reaches a blocking functions and waits
- ▶ Operating System tries to keep all run queues balanced and migrates threads
 - ▶ Good for most application, bad for parallel applications with frequent synchronization

Incrementing i 1/2

```
1  #define NUM 10000000
2
3  void * increment(void *i_void_ptr)
4  {
5      int *i = (int *) i_void_ptr;
6
7      for(int j=0; j < NUM; j++)
8          (*i)++;
9
10     return NULL;
11 }
```

Incrementing i 2/2

```
1 void main()  
2 {  
3     int i = 0;  
4     pthread_t thr;  
5     pthread_create(&thr, NULL, &increment, &i);  
6  
7     for(int j=0; j < NUM; j++)  
8         i++;  
9  
10    pthread_join(thr, NULL);  
11    printf("Value of i = %d\n", i);  
12 }
```

```
$ ./increment_integer
```

```
Value of i = 11315419
```

```
$ ./increment_integer
```

```
Value of i = 11038305
```

Data Hazards

Data hazards occur when threads are accessing shared data. Ignoring potential data hazards can result in a race condition. There are three situations in which a data hazard can occur.

- ▶ read after write (RAW), a "true dependency"
- ▶ write after read (WAR), an "anti-dependency"
- ▶ write after write (WAW), an "output dependency"

GCC Explorer Demo

- ▶ <http://gcc.godbolt.org/>

Incrementing i with Mutex 1/2

```
1  #define NUM 10000000
2
3  pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
4
5  void * increment(void *i_void_ptr)
6  {
7      int *i = (int *) i_void_ptr;
8
9      for(int j=0; j < NUM; j++)
10     {
11         pthread_mutex_lock(&mutex);
12         (*i)++;
13         pthread_mutex_unlock(&mutex);
14     }
15
16     return NULL;
17 }
```

Incrementing i with Mutex 2/2

```
1  void main()
2  {
3      int i = 0; pthread_t thr;
4
5      pthread_create(&thr, NULL, &increment, &i);
6
7      for(int j=0; j < NUM; j++)
8      {
9          pthread_mutex_lock(&mutex);
10         i++;
11         pthread_mutex_unlock(&mutex);
12     }
13
14     pthread_join(thr, NULL);
15
16     printf("Value of i = %d\n", i);
17 }
```

Incrementing i with Spinlock 1/2

```
1  #define NUM 10000000
2
3  pthread_spinlock_t spinlock;
4
5  void * increment(void *i_void_ptr)
6  {
7      int *i = (int *) i_void_ptr;
8
9      for(int j=0; j < NUM; j++)
10     {
11         pthread_spin_lock(&spinlock);
12         (*i)++;
13         pthread_spin_unlock(&spinlock);
14     }
15
16     return NULL;
17 }
```

Incrementing i with Spinlock 2/2

```
1  void main()
2  {
3      int i = 0; pthread_t thr;
4
5      pthread_spin_init(&spinlock, PTHREAD_PROCESS_PRIVATE);
6      pthread_create(&thr, NULL, &increment, &i);
7
8      for(int j=0; j < NUM; j++)
9      {
10         pthread_spin_lock(&spinlock);
11         i++;
12         pthread_spin_unlock(&spinlock);
13     }
14
15     pthread_join(thr, NULL);
16     printf("Value of i = %d\n", i);
17 }
```

Comparison Mutex and Spinlock

```
$ time ./increment_integer_mutex  
Value of i = 20000000
```

```
real    0m1.079s  
user    0m0.937s <-- user space  
sys     0m1.137s <-- kernel space
```

```
time ./increment_integer_spinlock  
Value of i = 20000000
```

```
real    0m1.062s  
user    0m2.067s <-- user space only  
sys     0m0.000s
```

Synchronization in General

- ▶ Bad for performance
- ▶ Serializes application (Amdahl's law)
- ▶ Hurts scalability
- ▶ Avoid if possible
 - ▶ Duplicate data
 - ▶ Rewrite algorithm
- ▶ Chose best synchronization primitive for your task

Assignment for this week: Dynamic Work Distribution

- ▶ In the last exercise on pi, the work was split in the beginning (statically) into number of threads pieces
- ▶ Each thread was computing his part of the work and in the end the result was combined
- ▶ This time the work should be distributed dynamically during runtime (STEPS)
- ▶ Protect shared variables with the introduced synchronization primitives and measure the runtime
- ▶ Compare the runtime of the static work distribution with the dynamic one.